

The Couchbase Blog

Couchbase, the NoSQL Database

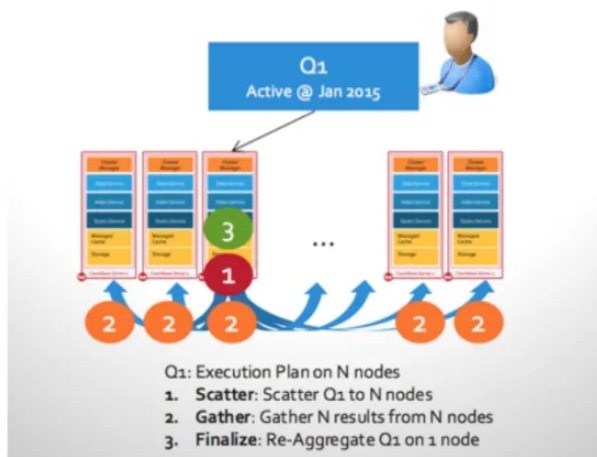
Comparing Couchbase Views with Couchbase N1QL & Indexing.



Keshav Murthy on December 4, 2017

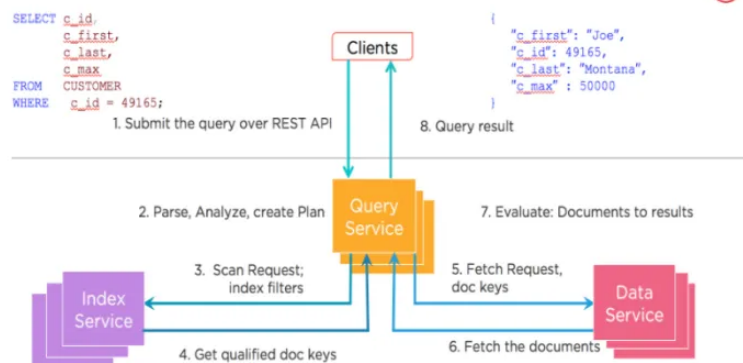
As Couchbase data platform evolved, services like N1QL and GSI Indexing handled the use cases Couchbase VIEWS used to handle and much more. It's logical to ask the comparative question between them. Here is a table comparing both. This is intended for developers and architects familiar with both them and not as an introductory article. Use the links here to learn more and play with the respective features.

Local Indexes (Views)



View Query Execution

N1QL: Query Execution Flow



N1QL Query Execution

Next →

Develop Web Applications w...

Topic

Couchbase Map-Reduce Views

Couchbase N1QL / GSI

Approach

Based on user-defined mapA

Based on declarative N1QL query / GSI

Approach	Based on user-defined map() and reduce() functions that operate on data in the background. Because map() and reduce() is written in Javascript, you can code complex logic within those functions.	Based on declarative N1QL query (SQL for JSON). Uses appropriate indexes to optimize execution and executed dynamically by orchestrating Query-Index-data services. N1QL enables easily writeable and readable queries for JSON. Because it's inspired by SQL, it's flexible, composable. Because, it's extended for JSON, it works on rich JSON data. Uses 4-valued boolean logic (true, false, NULL, MISSING)
----------	--	---

More Info	Couchbase Docs: http://bit.ly/2jQrY11	1. http://query.couchbase.com 2. https://blog.couchbase.com/n1ql-practical-guide-second-edition/
-----------	--	--

Querying	Query based on	Query Statements
	1. Single key	1. SELECT
	2. Set of keys	2. INSERT
	3. Start-End key	3. UPDATE
	4. Start-End document keys	4. DELETE
	5. Group BY, Aggregation	5. MERGE
	6. Pagination	6. INFER
		7. EXPLAIN

Query Operations:

Next →

Develop Web Applications w...

2. Set of keys

3. Range keys
4. Range of document keys
5. Arbitrarily complex predicates
6. INNER JOIN, LEFT OUTER JOIN
7. NEST, UNNEST
8. GROUP BY
9. Aggregation
10. Pagination (OFFSET, LIMIT)
11. Optimization
12. ORDER BY
13. HAVING
14. Subqueries (correlated, non-correlated)
15. Derived tables
16. SET operations: UNION, UNION ALL, EXCEPT, EXCEPT ALL, INTERSECT
17. Highly composable queries, meaning these operations can be simply combined with each other to express complex business questions and operations easily.

Indexing

Simple index for views.

1. Primary Index
2. Named primary index
3. Secondary index

Next →

Develop Web Applications w...

6. Array Index

- 7. ALL array
- 8. ALL DISTINCT array
- 9. Partial Index
- 10. Adaptive Index
- 11. Duplicate Indices
- 12. Covering Index

Partitioning	Aligned to data partitioning.	Independent services. N1QL and GSI scales independent of Data service and each other.
Scale	Scales with data service	Independent scaling via Multidimensional scaling (MDS)
Fetch with document key	Because the data is partitioned on document key, fetches the document directly from the node	Specify the query via USE KEYS clause. Because the data is partitioned on document key, fetches the document directly from the node
Fetch with Index key	Scatter-Gather	Each index scan on a single node; Data on multiple nodes. Post processing in Query node
Range scan	Scatter-Gather	Index scan on a single node.

Next →

Develop Web Applications w...

Grouping,	Built-in with Views API	Built into N1QL
-----------	-------------------------	-----------------

aggregation

Caching

File system

Index buffer pool

Data cache

Storage

Couchstore

Plasma storage engine (5.0 & above)

Memory Optimized Index (4.5 and above)

ForestDB (community)

Availability

Replica Based

5.0: Replicas

4.x: Equivalent Indexes

Query Latency

10 milliseconds to 100 milliseconds

5 milliseconds+

(Simple queries)

Query Throughput

3K to 4K queries per second

40K queries per second

(Simple queries)

Next →

Develop Web Applications w...

data service)

query services: MDS)

Applicability	Aggregations, best of large scale aggregations for low and moderate latency requirements. Map-reduce operations on the data is done in the background as the data is modified.	Best for attribute based lookup, range scans, complex select-join-project-array Operations. Supports grouping, aggregation and ordering — these operations are done dynamically as part of query execution.
Application requirements	Report on well defined metrics Large scale aggregations Latency sensitive	Secondary key lookups Range Scans Operational aggregations Filtered queries Ad-hoc queries with complex predicates, joins, aggregations, app search, pagination, secondary key based updates.
Spatial	Supported via Spatial Views	Not directly. https://dzone.com/articles/speed-up-spatial-search-in-couchbase-n1ql
Consistency	Stale = UPDATE_AFTER Stale = OK	Unbounded (stale = OK) AT_PLUS (read your own writes)

Next →

Develop Web Applications w...

updates up to now(). Stale = False).

Tools

Web console

Web console, Developer workbench,
Query monitoring, Query Profiling,
Visual explain, INFER.

Posted in: Application Design, Best Practices and Tutorials, Couchbase Architecture, Data Modeling, N1QL / Query

Tagged in: couchbase, Indexing, MapReduce, N1QL, performance, Secondary Indexing, SQL

Posted by Keshav Murthy

Keshav Murthy is a Vice President at Couchbase R&D. Previously, he was at MapR, IBM, Informix, Sybase, with more than 20 years of experience in database design & development. He lead the SQL and NoSQL R&D team at IBM Informix. He has received two President's Club awards at Couchbase, two Outstanding Technical Achievement Awards at IBM. Keshav has a bachelors degree in Computer Science and Engineering from the University of Mysore, India, holds eight US patents.

[All Posts](#) [Website](#)



Leave a reply

Comment

Next →

Develop Web Applications w...
